

Lists in R

Vectors (Again)

Recall that a **vector** is a sequence of entries of the *same data type*. The following are all vectors:

```
# Create numeric, character and logical vectors
numerics <- c(1, 2.5, 5, -1, 0.08, 0)
characters <- c("abc", "d", "6", "six", "l", "f")
logicals <- c(TRUE, FALSE, FALSE, TRUE, TRUE, TRUE)
```

At the risk of overemphasis, each of these 3 objects are vectors because the entries of each are of the same data type. The entries of the vector `numerics` are all numerical (real numbers or integers); the entries of `characters` are all characters; and the entries in `logicals` are all logical operators (TRUE/FALSE). To verify, you can check the Environment pane under Values where these vectors are listed with an abbreviated description of their data type (`chr`, `logi`, and `num`). You can also use the `class()` function, for example:

```
# Verify data type in characters vector
class(characters)
```

Lists

A data frame generalizes a vector in the sense that it combines vectors of the same length into its own data structure. Another way to generalize a vector is to allow different data types in the sequence of entries. A *list* is a sequence of entries that accommodates a mixture of data types (and much, much more). To create a list, instead of the `c()` function, use `list()` around its components:

```
# Create a list
my_list <- list(7, "Doom", FALSE, -98, TRUE, "ALL CAPS!")
```

All 3 data types are present in `my_list`, so it's certainly not a vector. Moreover, note that, upon running the command, R adds `my_list` under Data in the Environment pane. Clicking the blue-circle-with-white-arrow icon displays the data types of each of its contents.

As intimated in the (non-exhaustive) definition given above, lists are not limited to sequences of entries. List components may include data frames, matrices, results from statistical analyses, vectors, lists (and sublists of lists) and much more. The vast range of objects that are found in lists is why we use the word components, instead of entries. To illustrate the point, run these commands and examine the structure of `my_crazy_list` in the Environment pane:

```
# Create my_crazy_list
doom <- "ALL CAPS when you spell the man's name!"
num_vector <- c(12, 42, -12)
df <- data.frame(characters, logicals)
my_crazy_list <- list(doom, num_vector, df, list(doom, df))
my_crazy_list
```

It's complicated. To temper this type of complexity it's often wise to name a list's components, which then allows you to access components by name. Assigning names to components also clarifies the structure of a list and, consequently, makes it less likely you commit coding errors. Modify `my_crazy_list` to include names for its 4 list components:

```
# Add names to list components of my_crazy_list
my_crazy_list <- list(d = doom, n = num_vector, f = df, l = list(doom, df))
my_crazy_list
```

Accessing (Indexing) Elements of Lists

On account of their greater complexity, accessing components of lists is more complicated than subsetting vectors. To understand the component parts of a list, call the structure function, `str()`, as follows:

```
# Examine structure of my_list
str(my_list)
```

A key difference between indexing vectors and indexing lists concerns the use of the square brackets (`[]`). With a vector, calling, for example, `x[7]` returns the seventh element of the vector. But the command `my_list[4]` returns a list, and specifically a list containing one element `-98`. To extract the element `-98` and just the element, use double square brackets: `my_list[[4]]`. You should compare the results from the following commands to verify you understand the difference between the use of `[]` and `[[[]]` in lists.

```
# Accessing the 3rd and 4th elements of my_list as sublists
my_list[3]
my_list[4]
```

```
# Extracting the 3rd and 4th elements of my_list as entries
my_list[[3]]
my_list[[4]]
```

So, the double-square brackets, `[[[]]`, extract a list component and return it to its original data-type/data-structure. For example, if the list component is a vector, then using `[[[]]` will return a vector. There are in fact 3 ways to accomplish this. To verify, let's return to `my_crazy_list` and confirm that these 3 commands return the same object:

1. `my_crazy_list[[3]]`
2. `my_crazy_list[["f"]]`
3. `my_crazy_list$f`

Each of these is useful in different contexts. But remember that qualifying phrase: “return it to its original data-type/data-structure.” If you use single brackets, you'll always get a list back.

You can exclude an element by using a negative index. For example, `my_list[-5]` returns a sublist of all elements from `my_list` *except* the 5th element.

So What?

Lists are important R data structures because:

1. They allow you to organize and recall disparate information in a relatively simple way.
2. Results from many R functions return lists, and thus to work with these lists you'll need to know how to extract their components.

An important example to illustrate the second point is the `summary()` function, which returns, among other things, standard regressions results. For a regression model, `summary()` is a list of these results. Consider this contrived example and be sure to observe the output from the last command:

```
# Regressing mpg on weight
mtcars_reg <- lm(mpg ~ wt, data = mtcars)

# Get regression summary
summary(mtcars_reg)
```

```
# View structure of summary(mtcars_reg)
str(summary(mtcars_reg))
```

From that last command you should see a sublist of `summary(mtcars_reg)` called `coefficients` containing estimates as well as t - and p -values. You'll often need to extract coefficient estimates and their standard errors and, thus, you'll unavoidably have to apply the list rules from above. Let's look at the `coefficients` sublist:

```
# View estimates from mtcars_reg
summary(mtcars_reg)$coefficients
```

```
##              Estimate Std. Error   t value    Pr(>|t|)
## (Intercept) 37.285126   1.877627 19.857575 8.241799e-19
## wt          -5.344472   0.559101 -9.559044 1.293959e-10
```

To extract any entry in this 2×2 array, use square brackets, `[]`, with 2 indices within the square brackets: The first index referring to the row of the array, the second referring to the column of the array. If you've dealt with matrices before, this way to specify an entry of a matrix should sound familiar. That is, the entry in the i^{th} row and j^{th} column from matrix x is denoted by x_{ij} ; the row index preceding the column index. Likewise with lists: To extract the entry in the i^{th} row and j^{th} column of a list, you'll enter `[i,j]`. So you can think of square brackets as a type of matrix.

Let's consider some examples. To extract the slope estimate, `wt`, which is located in the 2nd row and the first column, run:

```
# Extract slope estimate
summary(mtcars_reg)$coefficients[2,1]
```

What about its standard error?

```
# Extract standard error of slope estimate
summary(mtcars_reg)$coefficients[2,2]
```

Dividing a coefficient's estimate by its standard error approximates its t -value. Let's verify for `wt`:

```
# Compute t-value of wt coefficient
summary(mtcars_reg)$coefficients[2,1]/summary(mtcars_reg)$coefficients[2,2]
```

This is a lot of typing with little payoff. Expedite the process by creating your own function that will extract estimates with much less effort.

The map Function

Create a list of 4 entries with each entry a set of randomly-generated numbers from different normal distributions:

```
set.seed(1325)

normal_list <- list(a = rnorm(n = 50, mean = 1, sd = 1),
                   b = rnorm(n = 10, mean = -0.5, sd = 0.5),
                   c = rnorm(n = 140, mean = 9, sd = 0.7),
                   d = rnorm(n = 7, mean = -7, sd = 6))
```

Note, for example, that list-element `c` contains 140 numbers, while `d` contains only 7 numbers.

If you wanted the mean and standard deviation from each of these list-elements you could call `mean()` and `sd()` 4 times apiece to obtain the summary statistics. But, creating a function is a much more efficient way to complete this task:

```
# Create function to compute mean and sd; return as a tibble
mean_and_sd <- function(x) {
  mean_x <- mean(x)
  sd_x <- sd(x)

  tibble(mean = mean_x,
         sd = sd_x)
}
```

Using `mean_and_sd()` determine the means and standard deviations of the 4 list components using appropriate list-extraction syntax (whichever of the 3 ways you prefer):

```
# Compute 4 means and sds from normal_list
mean_and_sd(normal_list[[1]])
```

```
## # A tibble: 1 x 2
##   mean    sd
##   <dbl> <dbl>
## 1  1.01 0.980
```

```
mean_and_sd(normal_list[[2]])
```

```
## # A tibble: 1 x 2
##   mean    sd
##   <dbl> <dbl>
## 1 -0.417 0.477
```

```
mean_and_sd(normal_list[[3]])
```

```
## # A tibble: 1 x 2
##   mean    sd
##   <dbl> <dbl>
## 1  8.98 0.742
```

```
mean_and_sd(normal_list[[4]])
```

```
## # A tibble: 1 x 2
##   mean    sd
##   <dbl> <dbl>
## 1 -6.80  5.82
```

I still had to copy-and-paste the first line 3 times and then replace the 1 with a 2 and then the 1 with a 3 and then the 1 with a 4. The `mean_and_sd()` function was more efficient than simply applying `mean()` and `sd()` to each of the lists. But there's still a better way to automate this task. How about a for loop?

As with all for loops, first allocate space to save the output with the `vector()` function. The only difference from past storage vectors is that, instead of each element of the vector being a double, save them as a list. In this case, you'll have 4 list components consisting of the 4 tibbles, each tibble containing a mean and standard deviation.

As for the body of the for loop, it's very similar to the last command except that instead of a numeric index, introduce the arbitrary index `i`.

```
# Create for loop to compute means and sds; save them in output
output <- vector("list", length(normal_list))

for (i in seq_along(normal_list)) {
  output[[i]] <- mean_and_sd(normal_list[[i]])
}
```

You should view `output` in the Environment tab to confirm it's in fact a list with the 4 mean-standard-deviation pairs as its list components.

This for loop proved to be the most efficient way to generate these means and standard deviations. But for loops aren't without flaws. For example, if you were reading the above code block, it'd take some time to figure out what's going on. The core of a for loop is its body. But the body is relatively small in comparison to the total amount of code required to define a for loop. Then there's always the book-keeping involved when performing iteration, which distracts from the for loops objective.

The `map()` function from the **purrr** package strips down iteration to its basics: It takes a vector/list/dataframe as input, applies a function to each of its pieces and returns a new object of the same length. Let's use `map()` to do what the for loop above does.

The first argument of `map()` is the vector/list/dataframe you want to iterate over. The second argument is the function you want applied to each piece. In our case, we'll apply the function `mean_and_sd()` to `normal_list`.

```
# Use map to return means and sds of normal_list
output_map <- map(normal_list, mean_and_sd)
```

Compare `output` to `output_map` to confirm the same object is returned, i.e., a list consisting of 4 tibbles, each with a mean and standard deviation.

The `map()` permits you to narrow your focus on the 2 core components of iteration: The object to be iterated over and the function applied to that object. All the iteration book-keeping and additional code necessary when crafting a for loop are behind the scenes when calling `map()`, making the code much clearer to comprehend.

Consider the following commands and see if you can explain to yourself what they do prior to running the code:

```
output_med <- map(normal_list, median)

output_iqr <- map(normal_list, IQR)
```